

# Project Proposal

Henry Gardner  
DATASCI 266 NLP

## Idea

Ever searched health symptoms online and were left more worried than before? I plan on building a model that will assess your described concerns and offer a discussion with suggestions and possible conditions, with an empathetic approach.

The main challenge is accurately classifying conditions. Currently, LLMs can offer some insight, but I seek to use the data listed below to fine-tune a few models to more accurately assess the said conditions. I am interested in exploring techniques such as LoRA for efficiency. From my experience, the “empathetic” responses can come through prompt engineering, although this might need to be further considered after the fine-tuning. Good implementations currently exist, in both open-source and LMaaS LLMs, however, I seek to improve upon LLM output for symptom screening. I am considering building out a custom algorithm for “model competition” allowing for multiple LLMs to compete for the output, or a RAG-based approach if I can find sufficient data, but this needs some more thought.

The goal is to give insight into medical applications to LLMs, not replace healthcare professionals. The importance of this lies in the speed and accessibility anyone could receive basic screenings. I love the idea of a free tool to help people discuss medical concerns, instead of being left worried after an internet search.

## Data

guozhuoran918. “GitHub - Guozhuoran918/NLICE.” GitHub, 2023, [github.com/guozhuoran918/NLICE](https://github.com/guozhuoran918/NLICE).

Mooney, Paul. “Medical Speech, Transcription, and Intent.” Kaggle.com, 2019, [www.kaggle.com/datasets/paultimothymooney/medical-speech-transcription-and-intent/data](https://www.kaggle.com/datasets/paultimothymooney/medical-speech-transcription-and-intent/data).

Shivade, Chaitanya. “MedNLI - A Natural Language Inference Dataset For The Clinical Domain” (version 1.0.0). PhysioNet (2019), <https://doi.org/10.13026/C2RS98>.

“Synthetichealth/Synthea.” GitHub, 6 Nov. 2020, [github.com/synthetichealth/synthea](https://github.com/synthetichealth/synthea).

## References

Brenna Li, Amy Wang, Patricia Strachan, Julie Anne Séguin, Sami Lachgar, Karyn C Schroeder, Mathias S Fleck, Renee Wong, Alan Karthikesalingam, Vivek Natarajan, Yossi Matias, Greg S Corrado, Dale Webster, Yun Liu, Naama Hammel, Rory Sayres, Christopher Semturs, and Mike Schaekermann. 2024. Conversational AI in health: Design considerations from a Wizard-of-Oz

dermatology case study with users, clinicians and a medical LLM. In Extended Abstracts of the 2024 CHI Conference on Human Factors in Computing Systems (CHI EA '24). Association for Computing Machinery, New York, NY, USA, Article 88, 1–10.

<https://doi.org/10.1145/3613905.3651891>.

Derek Tam, Margaret Li, Prateek Yadav, Rickard Bruel Gabrielsson, Jiacheng Zhu, Kristjan Greenewald, Mikhail Yurochkin, Mohit Bansal, Colin Raffel, & Leshem Choshen (2024). LLM Merging: Building LLMs Efficiently through Merging. In NeurIPS 2024 Competition Track, <https://openreview.net/forum?id=TiRQ4GI4lr>.

Ktena, I., Wiles, O., Albuquerque, I. *et al.* Generative models improve fairness of medical classifiers under distribution shifts. *Nat Med* 30, 1166–1173 (2024).

<https://doi.org/10.1038/s41591-024-02838-6>.

Singhal, Karan, et al. Towards Expert-Level Medical Question Answering with Large Language Models. 16 May 2023, <https://doi.org/10.48550/arxiv.2305.09617>.

Vaswani, Ashish Teku, et al. "Scale Efficiently: Insights from Pre-training and Fine-tuning Transformers." (2022),

<https://research.google/pubs/scale-efficiently-insights-from-pre-training-and-fine-tuning-transformers/>.